

Kontextsensitive Systeme

SoSe 2020

1 Introduction

- **Defining Context:**

- Parts of a discourse that surround a word and throw light on its meaning
- Synonym for: Surrounding, Situation, meta-information
- **Aspects:** Geographical, Physical, Organisational, Social, Emotional, User, Time, etc.
- Computer Science:
 - * Computer Linguistic: surrounding information of text
 - * System Architecture: virtual environment to suspend running program
 - * Lexical context which determines name resolution

- **Pervasive Computing Systems:**

- Convenient access through new class of appliances to relevant information
- with ability to take action on it **when and where you need**

- **Why context?**

- Human communication:
 - * situational info is implicit, to increase **conversational bandwidth**
- **Goal:** Make Interaction with computer systems more efficient
 - * limiting factor of IT: often not processor/memory but limits of human attention
 - * Mobile Computing: situation changes frequently and user is preoccupied with other tasks
 - * **Context-awareness is enabling technology** → overcome bottleneck

- **Different Aspects of context**

- Three important aspects: Schilit et. al
 - * Where are you? Who are you with? What resources are nearby?
- Four Dimensions of context: Gross and Specht
 - * Location, Identity (preferences, knowledge), Time (working hours), Environment or Activity

- **Context classification:**

- Active vs. Passive: set of environmental states and settings that
 - * active context: determines applications behaviour (critical)
 - * passive context: in which an application event occurs and is interesting to user (relevant)
- Primary vs. Secondary: information that characterizes situation of entity
 - * primary context: location - where, identity- who, time - when and activity - what
 - * secondary context: anything that can be found on basis of primary context
- Context-awareness of systems:
 - * if it uses context to provide relevant information / services to user
 - * where relevancy depends on users task

- **Classes of context-sensitive applications:**

1. Proximate selection & contextual information:

- **Trigger** = manual / explicit & **Output** = Information access
- **Examples:**
 - * google now → weather in the city (filtered and selected information)
 - * show nearest available printers based on location
 - * parking assistance that is triggered from driving backwards

2. Contextual commands:

- **Trigger** = manual / explicit & **Output** = Command execution
- **Examples:**
 - * Depending on the situation → emergency button call family / ambulance
 - * Use nearest printer for printing after command is executed
 - * multi purpose buttons that work differently in every context

3. Automatic contextual reconfiguration

- **Trigger** = automatic / implicit & **Output** = Information access
- **Examples:**
 - * change Notification from sound to vibration depending on situation
 - * use idle computers found nearby for storage rather than remote or local disks
 - * start parking assistance when driving forward very slowly

4. Context-triggered actions

- **Trigger** = automatic / implicit & **Output** = Command Execution
- **Example:**
 - * automatic closing / opening of windows depending on light conditions
 - * Contextual reminders like if at workplace show certain messages

- **Motivation: Context through Sensors:**

- Human Context Recognition:

- * Aristoteles: sight, hearing, touch balance, joint motion and acceleration
- * Additional: smell, taste, pain and thermoception

- Sensors:

- * Enable interaction by providing context
- * Number of smartphone sensors is increasing:
 - Light, Proximity, Acceleration, Touch, Rotation, Temperature, Compass, etc.
 - E.g. tilt → Orientation and Light-Sensor → Brightness

- **Classical Sensor Systems: Measurement Process**

- Classic measurement:

- * **Idea:** precisely control measurement conditions
- * Direct measurement, controlled placement and controlled environment
- * **But:** only appropriate if precision is more important than cost

- DIN 1319:

- * Define measurement task & unit as well as all conditions exactly
- * Select and calibrate measurement device and specify the process
- * Take the measurements and calculate noise influence and systematic error
- * Calculate overall result and quantitative error margin

- **Sensors and Context:**

- **Context** c transforms **signal** s to s^* at **state** x : $c(x, s) = s^*$
- Context Recognition: learn $c^{-1}(s, s^*) = \hat{x}$
 - * **But:** typically x is nominal, c is no function and c^{-1} does not exist

- **Complexity of Recognizing Contexts:**

- Complex context can only be sensed indirectly: affect environment and captured through sensors
- **Problem:** Sensors capture only partial effects and are noisy
- Representable Context:
 - * Information that can be encoded and is separable from activity
 - * Delineable: define contexts relevant for application in advance
 - * Stable: determination of relevance of contexts on an activity can be made once
- **Dey:** Context built-in at design time (context representable and processable)
- **Dourish:** technology becomes meaningful as individuals engage with it (context is emergent)

- **Detection Chain**

1. Physical Phenomenon: visible phenomenon in real world
2. Detection: Data comes from a sensor
3. Normalization: properly scale data
4. Feature Extraction & Reduction: compression, filtering and preparing of sensor data
5. Classification: interprets features and assigns class for them
6. Further abstraction: high level information and reasoning
7. Activity: action depending on output

- **Data Processing Chain:**

- Problem:
 - * Input: analog, continuous sensor signal
 - * Output: symbolic, discrete context class
- Sampling sensors yields discrete measurements
 - * High volume, depending on sampling rate
- Theory: sensor data as direct input for classifier
 - * But: greatly increases classifier complexity
 - * Solution extract information of importance

- **Resampling and windowing data**

1. Analog Sensor Signal:
 - Transform from time domain and specify period of time to be classified
 - **Problem:** Window lengths greatly affects classification results
2. Sensor Sampling:
 - Nyquist/Shannon not directly applicable: understanding theoretical possible
 - Reconstruction not necessary and features do not need to be perfect
3. Sample Windows:
 - Long: Smearing of context / latency (better: rolling window)
 - Short: Information does not characterize context and noisy results

2 Classification and Error

- **Step 1: Sample Windows**

- Specify period of time to be classified: segment data into windows of that time length
- **Problem:** Window length greatly affects classification results
- Practice: empirical evaluation for each application

- **Step 2: Feature Extraction**

- Goal: extract important information for context differentiation (classification) → reduce volume
- Importance of information is depends on context and sensors
- Features can be either in time or frequency domain (Fourier Transformation)

- **Step 3: Classification**

- Features append to feature vector → training with additional labels
- Classification: uses feature vector to predict next occurrence

- **K-Nearest Neighbour: most simple classifier**

- **Approach:** store all training vectors and labels
 - * For each classification vector: assign label based on "closest" vectors
 - * $d(p, q) = d(q, p) = \sqrt{\sum_{i=1}^n (q_i - p_i)^2}$
- Metrics for distance calculation
 - * Manhattan, Hamming, Mahalanobis, Euclidean
- **Good:** simple to implement, and search can be efficient through trees as storage
- **Bad:** low recognition rates, susceptible to normalization, lazy, memory consumption
- Big Data:
 - * Distances can be a problem, but if saving a lot of data is not an issue it can be quite good


- **Naive Bayes: Modelling by Probability**

- **Approach:** learn how data is distributed probabilistically and use it to infer context
- **Bayes Theorem:** $posterior = \frac{prior * likelihood}{evidence}$
- For Classification:
 - * Formula: $p(C|F_1, ..., F_n) = \frac{p(C)p(F_1, ..., F_n|C)}{p(F_1, ..., F_n)}$
 - * $\hat{c} = argmax_{k \in 1, ..., K} p(C_k) \prod_{i=1}^n p(F_i|C_k)$
- **Nomogram:** visualizing Naive Bayes - add logarithmic values
- **Good:** good results and low computational complexity
- **But:** assumption of uncorrelated data and oversimplified distributions

- **Decision Tree:**

- **Approach:** model data as tree such that nodes are decisions and leaves are labels
 - * Minimize statistical measures like Entropy or GINI and choose Split-Attribute with best value
 - * Repeat until datapoints are correctly classified
- Formulas:
 - * $Entropy(S) = -p_1 \log_2(p_1) - p_2 * \log_2(p_2)$ level of uncertainty
 - * $Gain(S, A) = Entropy(S) - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$


- Node impurity metrics for Decision Trees



- Example for calculations in Decision Trees:

Gini

Split on Class



Entropy

		Play Golf		
		Yes	No	
Outlook	Sunny	3	2	5
	Overcast	4	0	4
	Rainy	2	3	5
				14


$$\begin{aligned}
 E(\text{PlayGolf}, \text{Outlook}) &= P(\text{Sunny}) \cdot E(3,2) + P(\text{Overcast}) \cdot E(4,0) + P(\text{Rainy}) \cdot E(2,3) \\
 &= (5/14) \cdot 0.971 + (4/14) \cdot 0.0 + (5/14) \cdot 0.971 \\
 &= 0.693
 \end{aligned}$$

$$\begin{aligned}
 G(\text{PlayGolf}, \text{Outlook}) &= E(\text{PlayGolf}) - E(\text{PlayGolf}, \text{Outlook}) \\
 &= 0.940 - 0.693 = 0.247
 \end{aligned}$$

- Extensions:

- Dealing with numeric attributes:
 - discretise using a comparison: $<$, $>$, etc.
 - Computational expensive but: evaluate all possibilities and choose split with max Gain
- Further optimizations:
 - Random uniform subsampling and dynamic programming for reuse of results
- Reduced error pruning:
 - Replace leaf nodes with most popular class
 - Leave if classification is not affected
- Cost complexity pruning:
 - Remove subtree $prune(T, t)$ that minimizes: $\frac{err(prune(T, t), S) - err(T, S)}{|leaves(T)| - |leaves(prune(T, t))|}$

- Bias-Variance-Tradeoff: Model Complexity




- **Support Vector Machines:**

- **Approach:** create hyperplanes which separate data with highest margin (convex constraint opt.)
- Kernel trick: Increase dimensionality with transform \rightarrow linear classification in non-linear space
- **Good:** High accuracy and accuracy is not affected by high dimensions
- **Bad:** prone to overfitting, choice of kernel and parameters critical

- **Artificial Neural Networks**

- **Approach:** generate network of neurons (weights and thresholds) which minimizes a cost function
 1. Network with input for each feature and output for each class
 2. Network for each class with inputs and binary output for class
- **Neuron:** aggregation of weighted inputs with non-linear activation function
- **Training:** Backpropagation $\delta w_i = -\eta \frac{\partial E(w_i)}{w_i}$
- **Bad:** only broad heuristics for choosing the number of neurons, learnable through backpropagation

- **Classifier Overview:**



- **Dimensionality Reduction:**

- Autoencoder
 - * **Idea:** Input and Output-Layer have same size, hidden Layer has fewer neurons
 - * Lower dimensional representation through minimizing error between input and output
- Principal Component Analysis
 - * Highly covariant / correlated features deliver redundant information
 - * **Approach:**
 - Analyze covariance (correlation) matrix
 - Transform data into principle components (find eigenvectors and eigenvalues)
 - Ignore components with low weightings
 - * **Good:** removes covariance from features and reduces features and computation
 - * **Bad:** does not indicate value of a feature, all features must be calculated for transformation

- **How to find the right features:**

- Related work / empiric results / heuristics
- Brute force: try them all out and every combination
- Optimized: Feature selection algorithms

- **Machine Learning:**

- Unsupervised: PCA, Clustering, Autoencoder \rightarrow possible for feature learning
- Supervised: mapping inputs to outputs (classification / regression)
- Semi-Supervised: not every datapoint is labeled

3 Classification Error

- **Evaluating Classifiers: Training and Testing**

- **Goal:** Predict realistic error of classifier
- **But:** labels need to be known (supervised learning)
- Error Rate:
 - * Positives: instance is predicted correctly
 - * Negatives: instance is predicted incorrectly
 - * Accuracy: proportion of correct classifications over whole set for testing

- **Resubstitution Error Rate: Performance Test on Train-Data**

- Indicates only if code works and shows if function can be approximated
 - * does not estimate effect of unseen data
 - * not a good indicator for performance on future data
- Shows effect of model complexity but **error rate is not always 0%**
 - * Problems with data quality, statistical assumptions of the algorithm
 - * Features do not result in a clear discriminate classification
- **Solution:** Train-Test-Split
 1. Number of correct classifications: train error rate → lower is better
 2. Predictive Accuracy Evaluation: test error rate → lower is better

- **Training and Test-Set**

- **Idea:** independent samples that played no part in formation of testing rules
- **Assumption:** both samples are representative for underlying problem
- Sets may differ in context parameters: e.g. data from different countries
- **Operation Stages:**
 1. build basic structure
 2. optimize parameter settings, use (N:N) re-substitution but without test-data
- **Typically:** Train, Validation and Test data


- **Train-Test: Size comparison:**

- **Training:** larger set → better classifier
 - * **But:** error rate of Resubstitution (N:N) **only** tells if algorithm is applicable
- **Test:** larger set → more accurate error estimate
- **Dilemma:** how to choose the sizes if data is limited
 - * Holdout-method: one third for testing and rest for training
 - * Can be applied in repeated fashion and also include a validation set
 - * Once completed, all data can be used to build the final classifier

- **Confusion Matrix:**

Ground Truth \ Prediction	Prediction Positive	Prediction Negative	F ₁ -Score $2 * (PPV * R) / (PPV + R)$
	Condition Positive (P)	Condition Negative (N)	
Condition Positive (P)	True Positive (TP)	False Negative (FN)	Recall (R) TP / P
Condition Negative (N)	False Positive (FP)	True Negative (TN)	Specificity TN / N
Prevalence P / (P + N)	Precision (PPV) TP / (TP + FP)	NPV TN / (TN + FN)	Accuracy (TP+TN) / (PC+NC)

- **Receiver Operating Curve**



- **Singular metrics: better**

- F-Score: harmonic mean between precision and recall
- G-Score: geometric mean of precision and recall
- Cohens kappa: observed accuracy vs. expected accuracy

- **Repeated Holdout:**

- More reliable by repeating the process with sub-samples
 1. Each iteration selects proportion for random training, rest for testing
 2. Different error rates are averaged to get the overall error
- **Problem:** still no optimum, since the test sets overlap
- **Solution:** Cross-Validation (best 10-Fold CV)
 1. Split data into x subsets of equal size
 2. Use each subset in turn for testing and rest for training
 3. Error Rates are averaged to get the overall error

- **Leave-One-out CV (N-1;1)**

- **Idea:**

- * Divide data into set of m subsamples of equal size
 - * Use one sample for testing and train on remaining $(m-1)$ -samples
 - * This means the number of folds is equal to number of training instances

- **Computation:** build n classifiers for n instances

- **Error rate** = successful predictions / n

- **Good:** best use of data, no random subsampling,

- **Bad:** no stratification, expensive, assumption of statistical independence

- **Worst-Case-Example:** 2 classes with uniform distribution, prediction = Mode, Accuracy = 0

- **Dealing with statistical independence: e.g. time series**

- **Solution:** leave complete subject out of validation → measure inter-subject performance

- Often better represents real use case:

- * Learn classifier data from many users and apply for unknown user

- **Bootstrap: probability 0.632**

- **Idea:** better estimate for accuracy in small samples

- * Sample of N with N -times resampling as training-set
 - * Rest forms test set with probability $(1-1/n) = 0.368$

- Correction of test error: $err = 0.632 * err_{Test} + 0.368 err_{Train}$

- **Problem:** independence of both sets can be hard

- **Classifier Selection: No Free Lunch**

- ML is based on search and optimization of models and model parameters

- No Free Lunch Theorem:

- * for certain types of mathematical problems the cost of finding a solution
 - * averaged over all problems is the same for any solution method
 - * → no single best algorithm for all data

- But **which learning scheme performs better?** → 10-Fold CV for comparison

- **Model selection criterion:** MDL / Ocam's-Razor → smallest theory that describes all facts

- * Space required to describe theory + space required for mistakes

- **Improvements on CV:**

- Sample multiple times and check if mean accuracy for scheme "A" is better than "B"

- **In Practice:** limited data and limited estimates for computing the mean

- * paired- t-test: samples are paired and same CV is applied twice
 - * one-tailed binomial test: model accuracy better than no information
 - * McNemar-Test: check if model is biased, distribution same as in data

- **Classification with costs:**

- **Idea:** only predict high-cost class when very confident about prediction

- * Expected cost: dot product of vector of class probabilities and appropriate cost column
 - * Choose class that minimizes expected cost

- Cost-sensitive learning: resampling of instances / weighting of instances accordingly

- ML-Model may take costs into account like naive bayes

- **Ensemble Learning: Combining multiple models**

- **Idea:** build different simple experts and let them decide together
- **Good:** often improves predictive performance
- **Bad:** usually produces output that is hard to analyze (but there are solutions to some extent)

- **Bagging:**

- Combine predictions by voting/averaging: reduces variance and usually more classifiers are better
- Idealized version
 1. Build classifier for n splitted training sets
 2. Combine classifier predictions
- Learning scheme is unstable → small changes in data can have big impact
 - * Randomization: use random parameters or initial weights to build different classifiers
- **Using cost:** use confidence as weighting
- **Forest vs. SVM:** same solutions can be obtained

- **Boosting:**

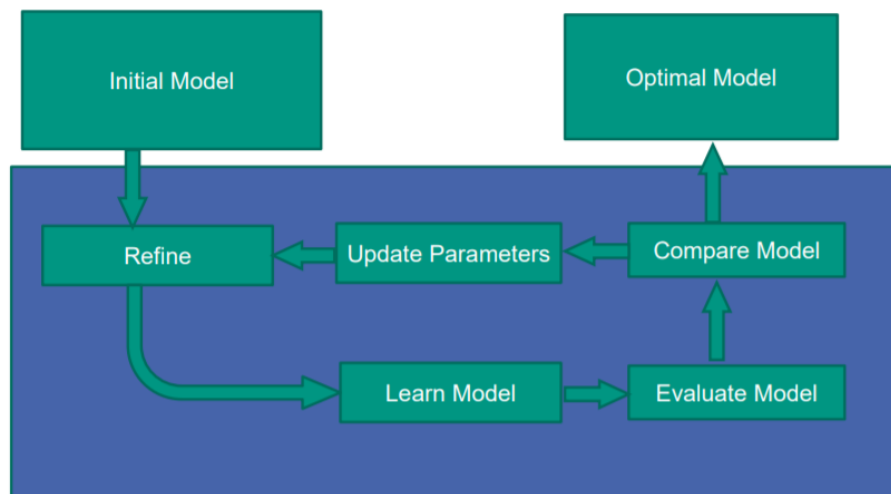
- Models are influenced by performance of previously built ones (also uses voting)
- Encourages new model to become expert for instances previously misclassified
- **Intuition:** combine weak learners to strong one
- **Forest vs. Adaboost:** similar but booster would perfectly separate if possible

- **Stacking: Stacked generalization**

- Combine predictions of base learners in meta learner (no voting)
- Usually different schemes are combined as input for meta learner
 - * CV-like: can't use predictions on training data to generate data for level 1 model
- Hard to analyze but can be used as initialization for automatic model selection

4 Optimizing Classification

- **Feature Selection:** Find best classifier by turning of features
 - Metaheuristics:
 - * Optimization-like techniques for feature subsets (sequentially turning on/off)
 - Filters:
 - * **Pipeline:** all features \rightarrow subset \rightarrow model \rightarrow performance
 - * Rank features using mutual information, correlation, similarity, etc.
 - * **Tradeoff** between: relevancy and redundancy: E.g. Minimum-Redundancy-Maximum-Relevance
 - * Select variables regardless of the model
 - * E.g. eliminate correlated features or least significant principle components
 - Wrappers:
 - * **Pipeline:** all features \rightarrow subset \rightleftharpoons model \rightarrow performance
 - * Evaluate subset of features using a learning a
 - Embedded Methods:
 - * **Pipeline:** all features \rightarrow subset \rightarrow model + performance
 - * Coupled with ML: e.g. recursive feature elimination of SVM (remove low weights)
 - * **Other methods:** L1 regularization (LASSO) or decision tree pruning
 - * Computational between Wrappers and Filters
- **General Requirements for Optimization Methods:**
 - Requirements: search method / blackbox but no gradient usage
 - Should include: flexible search space definition and scalable search (parallelize)
 - Maybe also: constraints
 - Blackbox-Optimization:



• Bayesian Optimisation

- **Idea:** use when function evaluation is expensive and gradient is not available
 - * Surrogate Model: approximate the Process
 - * Acquisition Function: is maximized balancing of exploitation (obj.) and exploration (surrogate)
- Choice of Surrogate Model:
 - * Many dimensions: Random Forest
 - * Categorical / hierarchical parameters: Gaussian process model with specialized kernel
 - * Otherwise: Gaussian process model with standard kernel
- Choice of Acquisition Function:
 - * Portfolio viable: use a portfolio of acquisition functions
 - * Absolute Min/Max known: MPI
 - * Otherwise: GP-LCB, EI, etc.
- **Good:** fully leveraged by choice of surrogate, data efficiency
- **Bad:** inherently sequential, output variable

• Evolutionary Algorithms:

- Initial Population: create initial population of random individuals
- Iteratively:
 1. Evaluation: compute objective values of solution candidates
 2. Fitness Assignment: use objective values to determine fitness values (ACC, R^2 , etc.)
 3. Selection: Select fittest individuals for reproduction (question about diversity)
 4. Reproduction: create new individuals from mating pool by crossover and mutation
- **TPOT-System:** Input Data + Cleaning but rest is done via evolutionary algorithms
- Neuro Evolution:
 - * Evolving instead of standard ML-Pipelines → built upon homogenous structures (Neurons)
 - * Good crossover-strategies exist but search space size increases gradually
 - * Small and thus complex models are preferred
 - * **Good:** prior knowledge initial population, parallelisation, output is variable
 - * **Bad:** data efficiency is bad due to much randomness

• Reinforcement-Learning:

- Agent tries to maximize his reward by taking actions which influence environment
- Usually modelled using MDPS:
 1. Controller samples architecture A with probability p
 2. Train child network with architecture A to get accuracy R
 3. Compute gradient of p and scale it by R to update controller
- **Good:** parallel, data efficient but converges slow, output is variable (e.g. sequence models)
- **Bad:** prior knowledge cannot be used

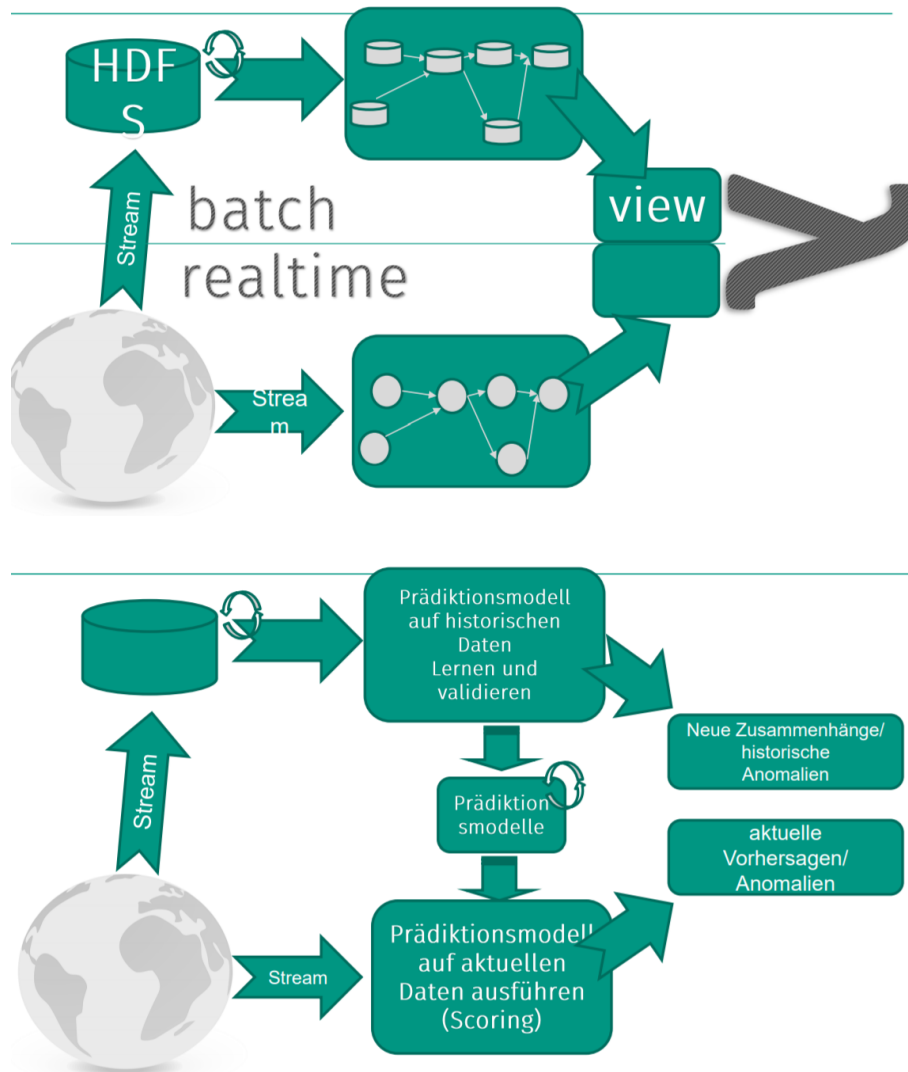
• Ensemble Learning & AutoML

- Classifiers can not only be used as surrogate function but also to find good initializations
- **Initializing:** find best model
 - * Build architecture similar to stacking
 - * Classifier that predicts how a classifier will perform and select good ones
 - * Maybe compare datasets: distribution, etc. and use it to predict accuracy of a classifier
 - * Use bag of best classifiers and construct ensemble → AutoML (Meta-Learning)

5 Big Data: Batch Processing

• Lambda-Architektur

- **Idea:** ML-Task cannot be started everytime from scratch → precalculation
 - * Query is split into realtime and batch tasks (e.g. CSS = offline learning + Realtime Classification)
 - * Batch-Layer: Hadoop or Spark (Answer bas)
 - * Realtime-Processing: Storm or Spark



• Realtime data processing: requirement for businesses

- Context-Sensitive Application: involves influence on decisions and latency / efficiency is important
- **Business Requirement:** reduce processing time and design strategy = zero-latency enterprise
- E.g. Context-aware online advertising, real-time search, high frequency trading, social networks
- Stream of events that flow into system at given time
 - * Complex Event Systems: typical "if than rules" and now use classifier every input
 - * E.g. Event-Pattern detection, event filtering, event aggregation and transformation
 - * Scoring: Prediction Model Markup Language (PMML), trained model is used in realtime system

- **PMML Support:**

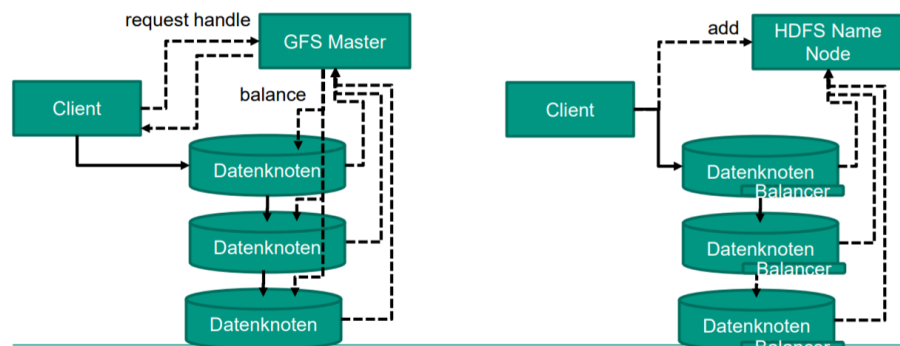
- SQL-Server: Decision Trees / Clustering Models (Distribution-based)
- Apache Spark: Clustering, K-Means, Regression Models (Linear, Ridge, Binary, Logistic), SVM
- PMML package: in R with lots of models and support for transformation
- **PMML Alternatives:** serialize as native structure (pickle RDS) and use same framework (sklearn)
- PFA (Portable Format for Analytics): successor to PMML

- **Hadoop: New way of batch processing**

- Software platform that lets one write and process vast amounts of data:
 - * MapReduce: offline computing engine
 - * HDFS: Hadoop distributed file system
 - * HBase: online data access
- Why is it useful:
 - * **Scalable:** reliable store and process petabytes
 - * **Economical:** distributes data and processing across clusters
 - * **Efficient:** distributing data and parallel processing on nodes where data is located
 - * **Reliable:** automatically maintains multiple copies and redeploys data if failure
- Assumptions: written with large clusters in mind
 - * Batch-Processing: emphasis on throughput opposed to low latency
 - * Large datasets: gigabyte to terabytes in HDFS
 - * Model: write-once-read-many access
 - * Moving: computation cheaper than moving data

- **Hadoop: What does it do?**

- Implements Googles MapReduce using HDFS which divides applications into many small blocks
- HDFS creates multiple replicas of data blocks for reliability
- MapReduce processes data where it is located (target Cluster with order of 10.000 nodes)



- **HDFS: Hadoop Distributed File System**

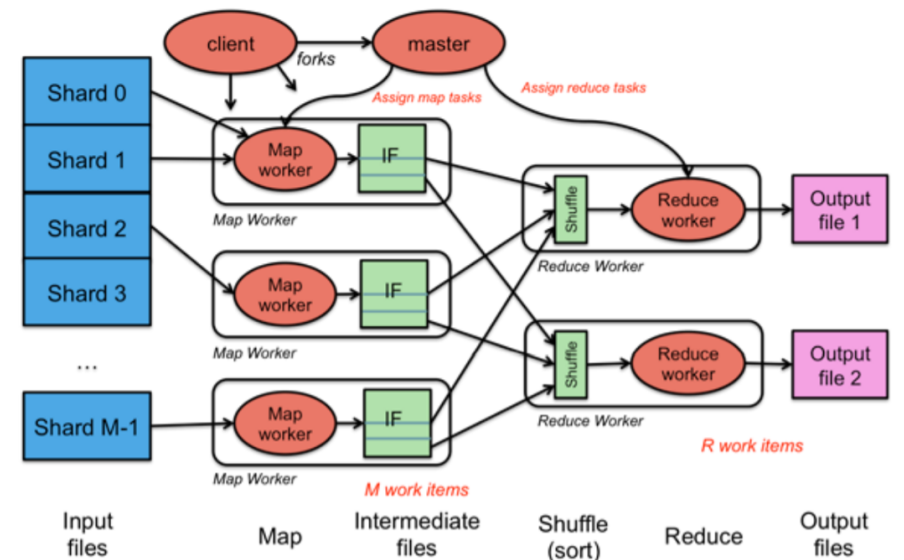
- File system designed to run on commodity hardware → similarities with existing systems
- **However:** differences to other distributed file systems are significant
 - * Highly fault-tolerant and designed to be deployed on low-cost hardware
 - * Provides high throughput access to application data and suitable for applications with large data
 - * relaxes POSIX requirements to enable streaming access to file system data
 - * Part of Apache Hadoop Core project

- **MapReduce Paradigm:**

- Intended as internal search/indexing application but now extensively used by more organizations
- Functional style programming with natural parallel capabilities across large clusters
- Map and Reduce are **user** written functions
- Underlying system takes care of:
 1. Partitioning input data and scheduling execution across several machines
 2. Handling machine failures and managing inter-machine communication (success key)
- Map-Phase: Datasets assigned to task tracker
 - * Run time partitions input and provides it to different Map-Instances
 - * Map $(key, value) \rightarrow (key^*, value^*)$
 - * Data functional operation will be performed emitting mapped key and value pairs (e.g. processing)
- Reduce phase:
 - * Run time collects $(key^*, value^*)$ pairs and distributes them to several Reduce function
 - * Each Reduce function gets pairs with the same key^* and produces single file output
 - * Master node collects answers to all subproblems and combines them in some way
 - * Forms answer to original problem (e.g. data collection and digesting, merging)

- **5 common steps of parallel computing:**

1. Prepare Map() input: row-wise and emit key value pairs per row \rightarrow Map input: $list(k1, v1)$
2. Run user-provided Map() code: \rightarrow Map output: $list(k2, v2)$
3. Shuffle Map output to Reduce processors, also group similar keys and input to same reducer
4. Run user-provided Reduce() code: custom reducer designed by dev and emit key and value
 - Reduce input: $(k2, list(v2))$
 - Reduce output: $(k3, v3)$
5. Produce final output: master node collects all reducer output and combines them in one file



- **R-tools: using Hadoop**

- **RHadoop:** Datanalytics with Hadoop via R functions
- **rhdfs:** R package providing Hadoop HDFS access to R, managing distributed files
- **rmr:** MapReduce interface for R

- **Challenges in Classification:**

- Models work on single in memory model: parallelization issue but only in learning phase
- High dimensional data: hadoop typically works by sorting in one dimension → transformation needed
- Hadoop build for batch processing:
 - * Realtime answer in classification phase can be a problem (e.g. waiting for one node)
 - * Mostly used for learning and score Models on different platform
 - * **Alternative:** streaming tricks → reduce chunk size / use custom scheduler

- **Comparison of different MapReduce Systems:**

- Classical hadoop: approach similar to read and write csv
 - * Data is sorted and location of data is known
 - * Good for big amounts of data but today in-memory is more efficient
- Spark: better utilizing in-memory
- Flink: allows for iterations until some state is reached
 - * Supports different workloads: stream, graph, iterative

- **KNN with Map-Reduce:**

- For amount of data approaching infinity → error rate $< 2 \times$ Bayes error rate
- Naive version: computation intensive for large sets → exact solution through linear scan or tree structure
- Scalable NN-Algorithm: e.g. through dimension reduction PCA
- Locality Sensitive Hashing:
 - * Hash-Function to map Datapoints that are near to each other
 - * Works in principle as dimension reduction
 - * Does not need to work perfect → for big k not all points need to be considered
 - * Typically Tradeoff between: accuracy and speed

- **Boosting with Map-Reduce:**

- Map-task: run boosting on its data and return sorted list of weak classifiers
- Reduce-Task: generate merged classifier and find its weight
- Key: use features as key and compute them on one node → good for sparse data
- XGBoost: Supported in many frameworks and almost linear in number of processors (trees)

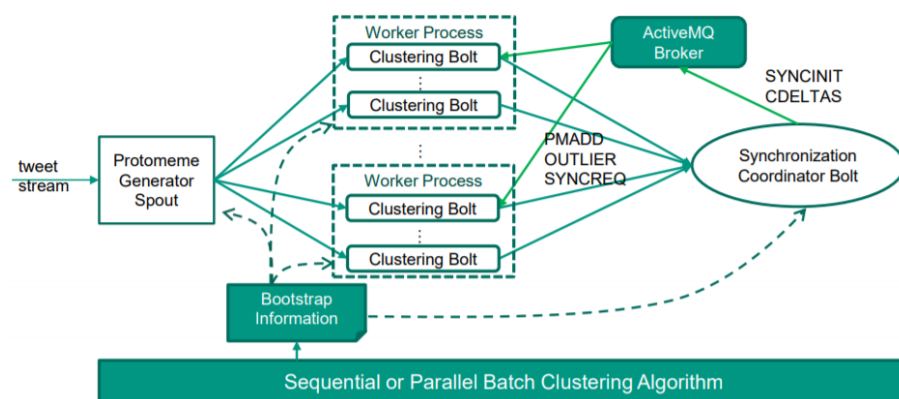
6 Realtime & Apps

• Microbatching vs. Streaming:

- Every stream can be represented as microbatch → small chunks of data
- **But:** keep data in-memory and avoid saving and reading from drive
- Why Real Time Stream Processing:
 - * Streams on Batch-System → often plagued by latencies
 - * Processing system must keep up with event rate → load shedding
 - * MapReduce / Hadoop store and process at scale but not for realtime systems → no hack viable
 - Latency can be reduced to get similar effects
 - But set of requirements differs fundamentally from batch processing


• Storm Concepts:

- **Streams:** unbounded sequence of tuples
- **Spout:** Source of Stream, e.g. read from Twitter streaming API
 - * Tuple Tree → Spout tuple not fully processed until all tuples in tree completed
 - * If not completed within specified timeout → tuple is replayed
 - * Reliability API → Guaranteeing Message Processing
- **Bolts:** Processes input streams and produces new streams, e.g. Functions, Filters, Aggregations
- **Topologies:** Network of spouts and bolts



• Kappa-Architectures

- Compared to Lambda-Architecture only streams and Log Files → no separation
- **Idea:** buffered data-streams from log files for analysis through Kafka



- **Why frameworks:**

- Multiple use cases: processing, computation, etc.
- Data types, size. velocity: scalability
- Mission critical data: fault-tolerance
- Time series / pattern analysis: reliability

- **Context adaption examples:**

- **Peer service:** take advantage of room projector for presentation
- **Location semantics:** disable phone ringtone in quiet places
- **Internals state:** decrease playback quality when battery power is low
- **User task:** show parking spots / gas stations when driving
- **Environmental conditions:** detailed indications when visibility is low

- **Problem: Mindset Mismatch**

- Software systems today are produced according to a manufacturing model
- **Mindset:** finished product is constructed and shipped and should act like any other machine
→ reliable but oblivious to surrounding
- Paradigmatic Shortcomings: traditional if / else
 - * **Software rigidity:** variability points are hard-coded in architecture, difficult adding new ones
 - * **Lack of modularity:** tight coupling of business and infrastructural code → difficult maintenance
 - * Mindset Mismatch: tools oblivious to context in which application runs → not adaptable software
 - * Can grow quite complex and same goes for software engineering pattern like strategies

- **Context-aware programming languages:**

- **Hypothesis:** lack of adaptability due to unavailability of context-aware programming languages
- **Solutions? AOP, FOP or COP** e.g. AOP with different aspects embedded in code
- Context Oriented Programming (COP):
 - * Have different Layers and selectively turn them on or off → Behavioural variation
 - * Instead of If-Else Structure: behaviour depends simply on context
 - * Design Concepts:
 - **Context group:** collection of environmental situations sharing same characteristics
 - **Context:** represents single environmental situation (connected / not-connected)
- Other approaches:
 - * Event Condition Action: trigger event if condition is matched and activate action (big and complex)
 - * Rapid-Prototyping: works with lots of visual tools and similar to Layering-Approach
 - * Logic Programming: Precondition (temp.) → Postcondition(modify temp) \wedge HTTP

- **Variability in Live Contexts: Outlook**

- Search for system that is: broad usable, adaptable and extensible, modular
- Up till now: no groundbreaking tools to get context-awareness available in systems

7 Multi-Sensor Activity Context Detection for Wearable Computing

- **Paper overview:**

- Wearable computing applications **central part** of user context is **human activity**
- It should be **automatically acquired** through **sensors** to avoid annoyance
- Activity is measured when it occurs through **sensors all over the body**
- This Paper deals with the these sensors → placement & extraction of data

- **System Architecture:**

- Acquisition system is modular designed and provides communication to host system
- Acceleration sensors are mounted on small boards which are wired → wireless connects PC
- Recognition Algorithm: Bayes Classification
 - * Naive Bayes classifier: $p(a|x) = \frac{p(a)}{p(x)} \prod_{i=1}^n p(x_i|a)$
 - * Features: running mean and variance computed over window of 50 sampels

- **Experimental Setup:**

- Goal: recognize everyday postures and activities: sitting, standing, walking, writing, hand shaking
- Sensor Placement and Number: all major joints (12 sensors)

- **Results and Discussion**

- Results get better the more sensors are used → activities with more sensors are better detected
- Physical activity is central for context-aware and user-centred applications
- Platform in this paper demonstrates context extraction using acceleration sensors
- Still work to do: sensors for more complex activities and inference with only acceleration sensors

8 Reducing the Dimensionality of Data with Neural Networks

- **Paper overview:**

- High-dimensional data can be converted to low-dimensional by MLP
- Idea: small central layer to reconstruct high-dimensional input vectors
- Fine-Tuning through gradient-descent and starting with good initial weights
- Can work much better than PCA for reduction of data

- **Dimensionality Reduction:**

- Facilitates classification, visualization and communication of high-dimensional data
- Most common used is PCA → find directions with greatest variance
- **Here:** Non-linear generalization of PCA with adaptive multilayer encoder Network
- Algorithm:
 - * Training a multi-layer auto-encoder can be problematic → Pre-Training to get good weights
 - * Restricted-Boltzmann machine can be used for training → finding good configurations
 - * After pretraining multiple layers of feature detectors, the model is "unfolded"